



Scalabitz

Efficient Collections Framework

What's a Blitz?

Blitz-chess is a style of
rapid chess play.



Blitz-chess is a style of
rapid chess play.

Knights have horses.





Horses run fast.

```
def mean(xs: Array[Float]): Float =  
  xs.par.reduce(_ + _) / xs.length
```

Is it a good idea to run `...par.map(` on large lists directly?



8

Let's say I have a somewhat large (several millions of items, or so) list of strings. Is it a good idea to run something like this:



2

```
val updatedList = myList.par.map(someAction).toList
```



Or would it be a better idea to group the list before running ...par.map(, like this:

```
val numberOfCores = Runtime.getRuntime.availableProcessors
val updatedList =
    myList.grouped(numberOfCores).toList.par.map(_.map(someAction)).toList.flatten
```

UPDATE: Given that `someAction` is quite expensive (comparing to `grouped`, `toList`, etc.)

[scala](#) [parallel-collections](#)

[share](#) | [edit](#) | [close](#) | [flag](#) | [protect](#)

edited Apr 7 '12 at 14:02

asked Apr 7 '12 at 13:51



Vilius Normantas
706 • 2 • 8 • 21

[add comment](#)

[start a bounty](#)

2 Answers

active

oldest

votes



11

Run `par.map` directly, as it already takes the number of cores into account. However, do not keep a `List`, as that requires a full copy to make into a parallel collection. Instead, use `Vector`.



[share](#) | [edit](#) | [flag](#)



[add comment](#)



7

As suggested, avoid using lists and `par`, since that entails copying the list into a collection that can be easily traversed in parallel. See the [Parallel Collections Overview](#) for an explanation.



As described in the [section on concrete parallel collection classes](#), a `ParVector` may be less efficient

tagged

[scala](#) × 20605

[parallel-collections](#) × 28

asked 1 year ago

viewed 283 times

active 1 year ago

CAREERS 2.0

Mobile UX Developer

SpotMe
Lausanne, Switzerland / ...

Engineering Manager

SpotMe
Lausanne, Switzerland / ...

Software Engineer (C++ / OpenGL)

Moka Studio
Lausanne, Switzerland

[More jobs near Lausanne...](#)

Related

55 [Preferred way to create a Scala list](#)

42 [Scala: Nil vs List\(\)](#)

7 [When a ConcurrentBag is better than a List?](#)

1 [ParVector map is not running in parallel](#)

4 [scala ranges versus lists performance on large](#)

Is it a good idea to run `...par.map(` on large lists directly?



Let's say I have a somewhat large (several millions of items, or so) list of strings. Is it a good idea to run something like this:

8

```
val updatedList = myList.par.map(someAction).toList
```




2


Or would it be a better idea to group the list before running ...par.map(, like this:

```
val numberOfCores = Runtime.getRuntime.availableProcessors
val updatedList =
    myList.grouped(numberOfCores).toList.par.map(_.map(someAction)).toList.flatten
```

UPDATE: Given that `someAction` is quite expensive (comparing to `grouped` , `toList` , etc.)

With Lists, operations
can only be executed
from left to right






Not your typical list.

Bon app.





Understanding parallel exists and find

7 I take a `List[Int]` and want to search for a value `x` where `x * 10 > 500` in parallel. So `exists` should return `true` if the list contains any value of 51 or greater.

```
def f(x: Int) = {
    println("calculating for " + x)
    Thread.sleep(100 - x)
    println("finished " + x)
    x * 10
}

val res = List.range(1, 100).par.exists(f(_) > 500)
```

Which gives results:

```
calculating for 1
calculating for 25
calculating for 50
calculating for 75
calculating for 13
finished 75          // <- first valid result found: 75 * 10 > 500
finished 50
calculating for 51  // but it kicks off more expensive calculations
finished 25
calculating for 26
finished 13
calculating for 14
finished 1
calculating for 2
finished 51
finished 26
calculating for 27  // and more
finished 14
calculating for 15
finished 2
calculating for 3
finished 27
calculating for 28
finished 15
```

tagged

scala × 20605

parallel-processing × 5138

parallel-collections × 28

asked 1 year ago

viewed 292 times

active 1 year ago

CAREERS 2.0

Software Engineer, Backend

SpotMe
Lausanne, Switzerland /...

Talented DevOps needed for exciting SaaS platform....

3S Mobile
Zurich, Switzerland / remote

(Senior) Java Developer (f/m)
HolidayCheck AG
Bottighofen, Switzerland /...

Linked

14 Scala Parallel Collections- How to return early?

6 Why doesn't scala's parallel sequences have a contains method?

Related

19 How do I replace the fork join pool for a Scala 2.9

Understanding parallel exists and find

7

I take a `List[Int]` and want to search for a value `x` where `x * 10 > 500` in parallel. So `exists` should return `true` if the list contains any value of 51 or greater.

1

```
def f(x: Int) = {  
    println("calculating for " + x)  
    Thread.sleep(100 - x)  
    println("finished " + x)  
    x * 10  
}  
  
val res = List.range(1, 100).par.exists(f(_) > 500)
```

```
def par: ParHashMap[A, B]
```

Returns a parallel implementation of this collection.

For most collection types, this method creates a new parallel collection by copying all the elements. For these collection, `par` takes linear time. Mutable collections in this category do not produce a mutable parallel collection that has the same underlying dataset, so changes in one collection will not be reflected in the other one.

Specific collections (e.g. `ParArray` or `mutable.ParHashMap`) override this default behaviour by creating a parallel collection which shares the same underlying dataset. For these collections, `par` takes constant or sublinear time.

All parallel collections return a reference to themselves.

returns a parallel implementation of this collection

Definition Classes [HashMap](#) → [CustomParallelizable](#) → [Parallelizable](#)



is apparently not enough

Apparently not enough

PARALLEL COLLECTIONS

Parallel Collection Conversions



Converting between sequential and parallel collections

Every sequential collection can be converted to its parallel variant using the `par` method. Certain sequential collections have a direct parallel counterpart. For these collections the conversion is efficient— it occurs in constant time, since both the sequential and the parallel collection have the same data-structural representation (one exception is mutable hash maps and hash sets which are slightly more expensive to convert the first time `par` is called, but subsequent invocations of `par` take constant time). It should be noted that for mutable collections, changes in the sequential collection are visible in its parallel counterpart if they share the underlying data-structure.

Sequential	Parallel
mutable	
Array	ParArray
HashMap	ParHashMap
HashSet	ParHashSet

Contents

- [Overview](#)
- [Concrete Parallel Collection Classes](#)
- [Parallel Collection Conversions](#)
 - [Converting between sequential and parallel collections](#)
 - [Converting between different collection types](#)
- [Concurrent Tries](#)
- [Architecture of the Parallel Collections Library](#)
- [Creating Custom Parallel Collections](#)
- [Configuring Parallel Collections](#)
- [Measuring Performance](#)

No amount of
documentation is
apparently enough



stackoverflow

Questions

Tags

Users

Badges

Unanswered

Can reduceLeft be executed in parallel?



I just started learning Scala, so please be patient :-)

5

I have a question about how reduceLeft behaves. Here an example:




```
List(1, 2, 3, 4, 5) reduceLeft (_ + _)
```

I wonder if the calculation can be done simultaneously, e.g.:

first round:

The `reduceLeft`
guarantees operations are
executed from left to right




Parallel and sequential collections sharing operations

There are several
problems here

How we see users





How users
see the docs



Bending the truth.

And sometimes we
were just slow



So, we have a new API now

```
def findDoe(names: Array[String]): Option[String] =  
{  
    names.toPar.find(_.endsWith("Doe"))  
}
```

Wait, you renamed a method?

```
def findDoe(names: Array[String]): Option[String] =  
{  
    names.toPar.find(_.endsWith("Doe"))  
}
```

Yeah, `par` already exists.
But, `toPar` is different.

```
def findDoe(names: Array[String]): Option[String] =  
{  
    names.toPar.find(_.endsWith("Doe"))  
}
```

```
def findDoe(names: Array[String]): Option[String] = {  
    names.toPar.find(_.endsWith("Doe"))  
}  
  
implicit class ParOps[Repr](val r: Repr) extends AnyVal {  
    def toPar = new Par(r)  
}
```

```
def findDoe(names: Array[String]): Option[String] = {  
    ParOps(names).toPar.find(_.endsWith("Doe"))  
}  
  
implicit class ParOps[Repr](val r: Repr) extends AnyVal {  
    def toPar = new Par(r)  
}
```

```
def findDoe(names: Array[String]): Option[String] = {  
    ParOps(names).toPar.find(_.endsWith("Doe"))  
}  
  
implicit class ParOps[Repr](val r: Repr) extends AnyVal {  
    def toPar = new Par(r)  
}  
  
class Par[Repr](r: Repr)
```

```
def findDoe(names: Array[String]): Option[String] = {  
  (new Par(names)).find(_.endsWith("Doe"))  
}  
  
implicit class ParOps[Repr](val r: Repr) extends AnyVal {  
  def toPar = new Par(r)  
}  
  
class Par[Repr](r: Repr)
```

```
def findDoe(names: Array[String]): Option[String] = {  
    (new Par(names)).find(_.endsWith("Doe"))  
}  
  
class Par[Repr](r: Repr)
```

But, `Par[Repr]` does not have the `find` method!

```
def findDoe(names: Array[String]): Option[String] = {  
  (new Par(names)).find(_.endsWith("Doe"))  
}  
  
class Par[Repr](r: Repr)
```

True, but `Par[Array[String]]`
does have a `find` method.

```
def findDoe(names: Array[String]): Option[String] = {  
  (new Par(names)).find(_.endsWith("Doe"))  
}  
  
class Par[Repr](r: Repr)  
  
implicit class ParArrayOps[T](pa: Par[Array[T]]) {  
  ...  
  def find(p: T => Boolean): Option[T]  
  ...  
}
```

More flexible!

More flexible!

- does not have to implement methods that make no sense in parallel

More flexible!

- does not have to implement methods that make no sense in parallel
- slow conversions explicit



No standard library collections were
hurt doing this.

More flexible!

- does not have to implement methods that make no sense in parallel
- slow conversions explicit
- non-intrusive addition to standard library

More flexible!


- does not have to implement methods that make no sense in parallel
- slow conversions explicit
- non-intrusive addition to standard library
- easy to add new methods and collections

More flexible!

- does not have to implement methods that make no sense in parallel
- slow conversions explicit
- non-intrusive addition to standard library
- easy to add new methods and collections
- import switches between implementations

```
def findDoe(names: Seq[String]): Option[String] = {  
    names.toPar.find(_.endsWith("Doe"))  
}
```

```
def findDoe(names: Seq[String]): Option[String] = {  
    names.toPar.find(_.endsWith("Doe"))  
}
```



```
def findDoe(names: Seq[String]): Option[String] = {  
    names.toPar.find(_.endsWith("Doe"))  
}
```

But how do I write generic code?

```
def findDoe[Repr[_]](names: Par[Repr[String]]) = {  
    names.toPar.find(_.endsWith("Doe"))  
}
```


```
def findDoe[Repr[_]](names: Par[Repr[String]]) = {  
    names.toPar.find(_.endsWith("Doe"))  
}
```

Par [Repr [String]] does not
have a find

```
def findDoe[Repr[_]: Ops](names: Par[Repr[String]]) = {  
    names.toPar.find(_.endsWith("Doe"))  
}
```

```
def findDoe[Repr[_]: Ops](names: Par[Repr[String]]) = {  
    names.toPar.find(_.endsWith("Doe"))  
}
```

We don't do this.

A black and white close-up photograph of Albert Einstein. He has his characteristic wild, grey hair and a full, bushy grey beard. His eyes are looking slightly to the right of the camera with a thoughtful expression. His hands are clasped together in front of him, resting on what appears to be the edge of a table or desk. The lighting is dramatic, coming from the side to highlight the contours of his face and the texture of his hands.

Make everything as simple as
possible, but not simpler.

```
def findDoe(names: Reducable[String])= {  
    names.find(_.endsWith("Doe"))  
}
```

```
def findDoe(names: Reducable[String]) = {  
    names.find(_.endsWith("Doe"))  
}
```

```
findDoe(Array(1, 2, 3).toPar)
```


```
def findDoe(names: Reducable[String])= {  
    names.find(_.endsWith("Doe"))  
}  
  
findDoe(toReducable(Array(1, 2, 3).toPar))
```

```
def findDoe(names: Reducable[String]) = {  
    names.find(_.endsWith("Doe"))  
}
```

```
findDoe(toReducable(Array(1, 2, 3).toPar))
```

```
def arrayIsReducable[T]: IsReducable[T] = { ... }
```

So let's write a program!



```
import scala.collection.par._

val pixels = new Array[Int] (wdt * hgt)
for (idx <- (0 until (wdt * hgt)).toPar) {

}
```

```
import scala.collection.par._

val pixels = new Array[Int] (wdt * hgt)
for (idx <- (0 until (wdt * hgt)).toPar) {
  val x = idx % wdt
  val y = idx / wdt
}
```

```
import scala.collection.par._

val pixels = new Array[Int] (wdt * hgt)
for (idx <- (0 until (wdt * hgt)).toPar) {
  val x = idx % wdt
  val y = idx / wdt
  pixels(idx) = computeColor(x, y)
}
```

```
import scala.collection.par._

val pixels = new Array[Int] (wdt * hgt)
for (idx <- (0 until (wdt * hgt)).toPar) {
  val x = idx % wdt
  val y = idx / wdt
  pixels(idx) = computeColor(x, y)
}
```

Scheduler not found!

```
import scala.collection.par._  
import Scheduler.Implicits.global  
  
val pixels = new Array[Int] (wdt * hgt)  
for (idx <- (0 until (wdt * hgt)).toPar) {  
    val x = idx % wdt  
    val y = idx / wdt  
    pixels(idx) = computeColor(x, y)  
}
```

```
import scala.collection.par._  
import Scheduler.Implicits.global  
  
val pixels = new Array[Int] (wdt * hgt)  
for (idx <- (0 until (wdt * hgt)) .toPar) {  
    val x = idx % wdt  
    val y = idx / wdt  
    pixels(idx) = computeColor(x, y)  
}
```

New parallel collections

33% faster!

Now


103 ms

Previously

148 ms


Workstealing tree scheduler

rocks!




Workstealing tree scheduler

rocks!



But, are there other interesting

Fine-grained uniform workloads are on the opposite side of the spectrum.





```
def mean(xs: Array[Float]): Float = {  
    val sum = xs.toPar.fold(0)(_ + _)  
    sum / xs.length  
}
```



```
def mean(xs: Array[Float]): Float = {  
    val sum = xs.toPar.fold(0)(_ + _)  
    sum / xs.length  
}
```

Now

15 ms

Previously

565 ms



But how?

```
def fold[T] (a: Iterable[T]) (z:T) (op: (T, T) => T) = {  
    var it = a.iterator  
    var acc = z  
    while (it.hasNext) {  
        acc = op(acc, it.next)  
    }  
    acc  
}
```

```
def fold[T](a: Iterable[T])(z:T)(op: (T, T) => T) = {  
    var it = a.iterator  
    var acc = z  
    while (it.hasNext) {  
        acc = box(op(acc, it.next))  
    }  
    acc  
}
```



```
def fold[T] (a: Iterable[T]) (z:T) (op: (T, T) => T) = {  
    var it = a.iterator  
    var acc = z  
    while (it.hasNext) {  
        acc = box(op(acc, it.next))  
    }  
    acc  
}
```

Generic methods cause boxing of primitives

```
def mean(xs: Array[Float]): Float = {  
    val sum = xs.toPar.fold(0) (_ + _)  
    sum / xs.length  
}
```

```
def mean(xs: Array[Float]): Float = {  
    val sum = xs.toPar.fold(0) (_ + _)  
    sum / xs.length  
}
```

Generic methods hurt performance
What can we do instead?

```
def mean(xs: Array[Float]): Float = {  
    val sum = xs.toPar.fold(0) (_ + _)  
    sum / xs.length  
}
```

Generic methods hurt performance
What can we do instead?

Inline method body!

```
def mean(xs: Array[Float]): Float = {
    val sum = {
        var it = xs.iterator
        var acc = 0
        while (it.hasNext) {
            acc = acc + it.next
        }
        acc
    }
    sum / xs.length
}
```

```
def mean(xs: Array[Float]): Float = {
    val sum = {
        var it = xs.iterator
        var acc = 0
        while (it.hasNext) {
            acc = acc + it.next
        }
        acc
    }
    sum / xs.length
}
```

Specific type
No boxing!
No memory allocation!

```
def mean(xs: Array[Float]): Float = {
    val sum = {
        var it = xs.iterator
        var acc = 0
        while (it.hasNext) {
            acc = acc + it.next
        }
        acc
    }
    sum / xs.length
}
```

Specific type
No boxing!
No memory allocation!

565 ms → 281 ms

2X speedup

```
def mean(xs: Array[Float]): Float = {
    val sum = {
        var it = xs.iterator
        var acc = 0
        while (it.hasNext) {
            acc = acc + it.next
        }
        acc
    }
    sum / xs.length
}
```

```
def mean(xs: Array[Float]): Float = {
  val sum = {
    var it = xs.iterator
    var acc = 0
    while (it.hasNext) {
      acc = acc + it.next
    }
    acc
  }
  sum / xs.length
}
```

Iterators? For Array?
We don't need them!

```
def mean(xs: Array[Float]): Float = {
    val sum = {
        var i = 0
        val until = xs.size
        var acc = 0
        while (i < until) {
            acc = acc + a(i)
            i = i + 1
        }
        acc
    }
    sum / xs.length
}
```

Use index-based access!

```
def mean(xs: Array[Float]): Float = {
    val sum = {
        var i = 0
        val until = xs.size
        var acc = 0
        while (i < until) {
            acc = acc + a(i)
            i = i + 1
        }
        acc
    }
    sum / xs.length
}
```

Use index-based access!

281 ms → 15 ms

19x speedup

Are those optimizations parallel-collections specific?

Are those optimizations parallel-collections specific?

No

Are those optimizations parallel-collections specific?

No

You can use them on sequential collections

```
def mean(xs: Array[Float]): Float = {  
    val sum = xs.fold(0)(_ + _)  
    sum / xs.length  
}
```

```
import scala.collections.optimizer._  
def mean(xs: Array[Float]): Float = optimize{  
    val sum = xs.fold(0) (_ + _)  
    sum / xs.length  
}
```

```
import scala.collections.optimizer._  
def mean(xs: Array[Float]): Float = optimize{  
    val sum = xs.fold(0) (_ + _)  
    sum / xs.length  
}
```

You get 38 times speedup!

Future work



@specialized collections

- Maps
- Sets
- Lists
- Vectors



Both faster &
consuming less
memory

@specialized collections

- Maps
 - Sets
 - Lists
 - Vectors
- 
- Both faster & consuming less memory
- Expect to get this for free inside optimize{} block

jdk8-style streams(parallel views)

- Fast
 - Lightweight
 - Expressive API
 - Optimized
- 
- Lazy data-parallel operations made easy

Future's based asynchronous API

```
val sum = future{ xs.sum }
val normalized = sum.andThen(sum => sum/xs.size)
```

Boilerplate code, ugly

Future's based asynchronous API

```
val sum = xs.toFuture.sum  
val scaled = xs.map(_ / sum)
```

- Simple to use
- Lightweight
- Expressive API
- Optimized



Asynchronous data parallel operations made easy

Current research: operation fusion

```
val minMaleAge = people.filter(_.isMale)
    .map(_.age).min
val minFemaleAge = people.filter(_.isFemale)
    .map(_.age).min
```

Current research: operation fusion

```
val minMaleAge = people.filter(_.isMale)
    .map(_.age).min
val minFemaleAge = people.filter(_.isFemale)
    .map(_.age).min
```

- Requires up to 3 times more memory than original collection
- Requires 6 traversals of collections

Current research: operation fusion

```
val minMaleAge    = people.filter(_.isMale)
                           .map(_.age).min
val minFemaleAge = people.filter(_.isFemale)
                           .map(_.age).min
```

- Requires up to 3 times more memory than original collection
- Requires 6 traversals of collections

We aim to reduce this to single traversal with no additional memory.

Without you changing your code